

Understanding Constraints in Subqueries and Views



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Use subqueries within Hive queries

Understand the limitations of subqueries in Hive

Create and work with views in Hive

Set Operations in Hive

Set Operations



Union



Minus



Intersect

Set Operations



Union



Minus



Intersect

Union and Union All



Union

**Combine results
from 2 queries
with the same
column types**

Union and Union All



Union

Union removes duplicates


Union All preserves duplicates

Demo

**Set operations on queries using union
and union all**


Subqueries in Hive

Subqueries




Queries are run on tables

Subqueries



**Queries are run on results
in a tabular format**

Tabular Format




Tables created using the **create table** command

The **result set** of a select statement

A subquery

Tabular Format




Subqueries in Hive can be used in the:

- **FROM** clause
- **WHERE** clause

Demo

**Use subqueries in the FROM clause of
Hive queries**

The Where Clause



Supports two types of subqueries

- IN/NOT IN
- EXISTS/NOT EXISTS

```
select id from customers  
where id in (1111, 3333, 5555);
```

A Typical Query with an **IN** Statement

A single list of values to check whether they match with the column values


```
select id from customers
where id in
(select customer_id from orders);
```

A Query with an **IN** Statement

The subquery should be for a **single** column value


The subquery should return a **list** of column values

The subquery **should not reference** the parent query

Demo

**WHERE - IN and WHERE - NOT IN
subqueries in Hive**

```
select id from customers
where exists (
  select customer_id from orders
where orders.customer_id = customers.id
);
```



A Query with an **EXISTS** Statement

The subquery **should be correlated** i.e. reference the parent query

Demo


**WHERE - EXISTS and WHERE - NOT
EXISTS subqueries in Hive**

Demo

Create table and insert data using a subquery

Views in Hive

Views



View

**A virtual table with a subset of data
from a larger, more extensive table**

```
create view product_quantities  
as
```

```
select product_id, quantity  
from orders
```

Views with a Subset of Data

Only the result set of this query is available in the view

Views



Stored as a query in Hive's **metastore**

Executed when **used**

Updated when data in the **underlying table changes**

Contains data from **single or multiple tables**


Frozen in time, **not affected** by table changes

Demo

Create a view

Operations that are permitted on a view

Advantages of Views



Reduce query complexity

Restrict access to data

Construct different logical tables from the same physical table

Demo

Create multiple views of the same table

Alter a view

Overview

Use subqueries and understand its limitations in Hive

Use views to simplify subqueries and restrict access to data